

In [2]: `# Tom Deng 662007936`

```
In [3]: import numpy as np
import matplotlib.pyplot as plt
from skimage.io import imread
from skimage.transform import resize
from sklearn.cluster import KMeans

# Step 1: Load the image
image_path = "C:\\Users\\Tom\\Downloads\\seg2.jpg" # Change this to the actual path
image = imread(image_path)

# Step 2: Resize image for faster processing (optional)
image = resize(image, (300, 400), anti_aliasing=True) # Resize to 300x400 pixels

# Step 3: Reshape image to 2D array (each pixel is a data point)
h, w, c = image.shape # Get height, width, and channels
pixels = image.reshape(-1, c) # Reshape to (num_pixels, num_channels)

# Step 4: Apply K-Means clustering
num_clusters = 6 # Number of segments (adjust as needed)
kmeans = KMeans(n_clusters=num_clusters, random_state=42, n_init=10)
labels = kmeans.fit_predict(pixels)

# Step 5: Replace pixel values with cluster centers
segmented_image = kmeans.cluster_centers_[labels]
segmented_image = segmented_image.reshape(h, w, c) # Reshape back to original dimensions

# Step 6: Display the original and segmented images
fig, ax = plt.subplots(1, 2, figsize=(12, 6))
ax[0].imshow(image)
ax[0].set_title("Original Image")
ax[0].axis("off")

ax[1].imshow(segmented_image)
ax[1].set_title("Segmented Image (K-Means)")
ax[1].axis("off")

plt.show()
```

Original Image



Segmented Image (K-Means)



```
In [7]: #Problem 2
import pandas as pd
# Load the dataset (Modify 'housing_prices.csv' to the actual file name)
data = pd.read_txt("C:\\Users\\Tom\\Downloads\\housing_prices.txt") # Assum

# Extract features (X) and target variable (y)
X = data.iloc[:, 0].values.reshape(-1, 1) # Population
y = data.iloc[:, 1].values.reshape(-1, 1) # Price

# Normalize features (Mean Normalization)
X_mean = np.mean(X)
X_std = np.std(X)
X = (X - X_mean) / X_std

# Add intercept term (bias)
X = np.c_[np.ones((X.shape[0], 1)), X] # Add a column of ones to X
```

```
-----
AttributeError                                Traceback (most recent call last)
Cell In[7], line 4
      2 import pandas as pd
      3 # Load the dataset (Modify 'housing_prices.csv' to the actual file name)
----> 4 data = pd.read_txt("C:\\Users\\Tom\\Downloads\\housing_prices.txt")
# Assumes two columns: Population, Price
      6 # Extract features (X) and target variable (y)
      7 X = data.iloc[:, 0].values.reshape(-1, 1) # Population

AttributeError: module 'pandas' has no attribute 'read_txt'
```

```
In [8]: def compute_cost(X, y, theta):
        """
        Compute the cost function J for linear regression.

        Parameters:
        X (ndarray): Feature matrix (with bias term).
        y (ndarray): Target variable.
        theta (ndarray): Model parameters.

        Returns:
        float: Cost J
        """
        m = len(y)
        predictions = X.dot(theta)
        cost = (1 / (2 * m)) * np.sum((predictions - y) ** 2)
        return cost

def mini_batch_gradient_descent(X, y, theta, learning_rate, batch_size, num_
    """
    Perform Mini-Batch Gradient Descent.

    Parameters:
    X (ndarray): Feature matrix.
    y (ndarray): Target variable.
```

```

theta (ndarray): Model parameters.
learning_rate (float): Learning rate.
batch_size (int): Mini-batch size.
num_iters (int): Number of iterations.

Returns:
theta (ndarray): Optimized parameters.
J_history (list): History of cost function values.
"""
m = len(y)
J_history = []

for i in range(num_iters):
    indices = np.random.permutation(m) # Shuffle data indices
    X_shuffled = X[indices]
    y_shuffled = y[indices]

    for j in range(0, m, batch_size):
        X_batch = X_shuffled[j:j+batch_size]
        y_batch = y_shuffled[j:j+batch_size]

        predictions = X_batch.dot(theta)
        errors = predictions - y_batch
        gradient = (1 / batch_size) * X_batch.T.dot(errors)
        theta -= learning_rate * gradient # Update theta

    J_history.append(compute_cost(X, y, theta)) # Store cost for plotting

return theta, J_history

```

```

In [9]: # Initialize parameters
theta = np.zeros((2, 1)) # Two parameters: intercept and slope
learning_rate = 0.01
num_iters = 1000
batch_sizes = [1, 5, 10, 20]

# Store cost history for each batch size
cost_histories = {}

plt.figure(figsize=(10, 6))

for batch_size in batch_sizes:
    theta_init = np.zeros((2, 1)) # Reset theta for each batch size
    theta_opt, J_history = mini_batch_gradient_descent(X, y, theta_init, learning_rate, num_iters)
    cost_histories[batch_size] = J_history
    plt.plot(range(len(J_history)), J_history, label=f'Batch Size {batch_size}')

plt.xlabel("Iterations")
plt.ylabel("Cost J")
plt.title("Cost Function J over Iterations for Different Batch Sizes")
plt.legend()
plt.show()

```

```

-----
NameError                                Traceback (most recent call last)
Cell In[9], line 14
     12 for batch_size in batch_sizes:
     13     theta_init = np.zeros((2, 1)) # Reset theta for each batch size
--> 14     theta_opt, J_history = mini_batch_gradient_descent(X, y, theta_i
nit, learning_rate, batch_size, num_iters)
     15     cost_histories[batch_size] = J_history
     16     plt.plot(range(len(J_history)), J_history, label=f'Batch Size {b
atch_size}')

NameError: name 'X' is not defined
<Figure size 1000x600 with 0 Axes>

```

```

In [10]: # Normalize the input population value
pop_input = (160000 - X_mean) / X_std
pop_input = np.array([1, pop_input]).reshape(1, -1) # Add bias term

# Predict using the best-trained model (batch size 10 as an example)
predicted_price = pop_input.dot(theta_opt)
print(f"Predicted house price for a city with population 160,000: ${predicted_price}")

```

```

-----
NameError                                Traceback (most recent call last)
Cell In[10], line 2
     1 # Normalize the input population value
----> 2 pop_input = (160000 - X_mean) / X_std
     3 pop_input = np.array([1, pop_input]).reshape(1, -1) # Add bias term
     5 # Predict using the best-trained model (batch size 10 as an example)

NameError: name 'X_mean' is not defined

```

```

In [11]: #Problem 3
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.feature_selection import RFE
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

```

```

In [12]: # Load the breast cancer dataset
data = load_breast_cancer()
X = pd.DataFrame(data.data, columns=data.feature_names) # Feature matrix
y = data.target # Target variable (0 = Malignant, 1 = Benign)

# Display dataset information
print("Dataset Features:\n", X.head())
print("\nTarget Labels:\n", np.unique(y, return_counts=True)) # Distribution of target labels

```

Dataset Features:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	\
0	17.99	10.38	122.80	1001.0	0.11840	
1	20.57	17.77	132.90	1326.0	0.08474	
2	19.69	21.25	130.00	1203.0	0.10960	
3	11.42	20.38	77.58	386.1	0.14250	
4	20.29	14.34	135.10	1297.0	0.10030	

	mean compactness	mean concavity	mean concave points	mean symmetry	\
0	0.27760	0.3001	0.14710	0.2419	
1	0.07864	0.0869	0.07017	0.1812	
2	0.15990	0.1974	0.12790	0.2069	
3	0.28390	0.2414	0.10520	0.2597	
4	0.13280	0.1980	0.10430	0.1809	

	mean fractal dimension	...	worst radius	worst texture	worst perimeter	\
0	0.07871	...	25.38	17.33	184.60	
1	0.05667	...	24.99	23.41	158.80	
2	0.05999	...	23.57	25.53	152.50	
3	0.09744	...	14.91	26.50	98.87	
4	0.05883	...	22.54	16.67	152.20	

	worst area	worst smoothness	worst compactness	worst concavity	\
0	2019.0	0.1622	0.6656	0.7119	
1	1956.0	0.1238	0.1866	0.2416	
2	1709.0	0.1444	0.4245	0.4504	
3	567.7	0.2098	0.8663	0.6869	
4	1575.0	0.1374	0.2050	0.4000	

	worst concave points	worst symmetry	worst fractal dimension
0	0.2654	0.4601	0.11890
1	0.1860	0.2750	0.08902
2	0.2430	0.3613	0.08758
3	0.2575	0.6638	0.17300
4	0.1625	0.2364	0.07678

[5 rows x 30 columns]

Target Labels:

```
(array([0, 1]), array([212, 357], dtype=int64))
```

```
In [13]: # Split data (70% training, 30% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, ran
```

```
In [14]: # Standardize the features (zero mean, unit variance)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [15]: # Initialize logistic regression model
log_reg = LogisticRegression(max_iter=1000, random_state=42)

# Use RFE to select the top 2 features
rfe = RFE(log_reg, n_features_to_select=2)
```

```
rfe.fit(X_train_scaled, y_train)

# Get the selected features
selected_features = X.columns[rfe.support_]
print("\nSelected Best Two Features for Classification:", selected_features)
```

Selected Best Two Features for Classification: Index(['worst area', 'worst concave points'], dtype='object')

```
In [16]: # Select only the best two features
X_train_selected = X_train_scaled[:, rfe.support_]
X_test_selected = X_test_scaled[:, rfe.support_]

# Train logistic regression with selected features
log_reg.fit(X_train_selected, y_train)
```

```
Out[16]: LogisticRegression
LogisticRegression(max_iter=1000, random_state=42)
```

```
In [17]: # Predict on test set
y_pred = log_reg.predict(X_test_selected)
```

```
In [18]: # Compute classification metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)

# Display results
print("\nModel Evaluation Metrics:")
print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1 Score: {f1:.4f}")

# Display confusion matrix
print("\nConfusion Matrix:")
print(conf_matrix)

# Classification report
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

#### Model Evaluation Metrics:

Accuracy: 0.9415  
Precision: 0.9450  
Recall: 0.9626  
F1 Score: 0.9537


#### Confusion Matrix:

```
[[ 58   6]
 [   4 103]]
```

#### Classification Report:

	precision	recall	f1-score	support
0	0.94	0.91	0.92	64
1	0.94	0.96	0.95	107
accuracy			0.94	171
macro avg	0.94	0.93	0.94	171
weighted avg	0.94	0.94	0.94	171

```
In [19]: # Plot confusion matrix
plt.figure(figsize=(6, 4))
plt.imshow(conf_matrix, cmap="Blues", interpolation="nearest")
plt.colorbar()
plt.xticks([0, 1], ['Malignant', 'Benign'])
plt.yticks([0, 1], ['Malignant', 'Benign'])
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix")
plt.show()
```



In [20]:

In [ ]: